



April 9, 2024

TO: Members, Assembly Privacy and Consumer Protection Committee

**SUBJECT: AB 2930 (BAUER-KAHAN) AUTOMATED DECISION TOOLS
OPPOSE – AS INTRODUCED FEBRUARY 15, 2024
SCHEDULED FOR HEARING – APRIL 16, 2024**

The undersigned organizations must respectfully **OPPOSE AB 2930 (Bauer-Kahan)** as introduced February 15, 2024, relating to the development and deployment of automated decision tools (ADTs). We agree that it is critical that companies take care to reduce bias and discrimination in consequential decisions impacting people. At the same time, we believe these issues exist whether these decisions are human made from start to finish, or a biproduct of using or incorporating new technologies in the decision-making process. As such, it is critical that any regulatory efforts proceed with caution and precision, particularly as technology is still developing and has the potential to reduce, if not one day eliminate, such undesirable outcomes. There are many other beneficial uses of this technology as well. Just to name a few: they can enable quick approvals and access to credit that would otherwise take much longer, provider broader access to credit, protect consumers against fraudsters by assisting in the identification of uncharacteristic account activity, automate repetitive tasks (such as entering data in two places at once) and minimize errors by comparing current work to past work,

While we remain concerned about the sweeping scope of this bill both in terms of the types of ADT captured and industries impacted by **AB 2930**, we acknowledge and appreciate many of the critical changes the author has made to this legislation since the introduction of this concept in AB 331, last year. Recognizing that this list is not comprehensive (e.g., members have concerns stemming from the definitions of “developers” and “deployers”), our primary concerns largely center around the following issues, which our letter will further elaborate on. below:

- Scoping problems. The bill applies to every industry imaginable and is not sufficiently limited to high-risk uses of ADT within those industries. Both the scope and various key terms (in particular consequential decisions, automated decision tools, algorithmic discrimination, and controlling factor), all require additional clarity and/or narrowing.
- Unworkable obligations in relation to the mandated impact assessments and notice requirements. In particular, the “opt-out” related requirements, including the requirement that businesses allow consumers to opt out of an ADT when “technically feasible”.
- Overly narrow confidentiality protections. Impact assessments must be exempt from the California Public Records Act to encourage candor and avoid concerns around assessments becoming fodder for litigation.

- Allowing multiple enforcement entities will invariably create confusion for compliance. Providing for a single enforcer (the Attorney General) will promote consistent interpretation and application across the state.
- Insufficient preemption protections and related concerns around regulatory activity. Preemption is needed both to prevent conflict with localities and, similarly, to prevent concerns around departments/agencies over-regulating this technology and getting ahead of the Legislature and Governor. These issues are too important to Californians across the state and our struggling economy to significantly delegate and defer to unelected officials.

AB 2930's scope is both inordinately broad and vague, and likely to impede efforts to actually reduce bias and discrimination as a result of overregulation.

We understand the primary goal of this bill is to ensure responsible development and deployment of ADT to prevent algorithmic discrimination—a goal, once again, that we share. To do that however, we need clarity as to what is and is not considered discrimination. To be clear, we believe that algorithmic discrimination, or discrimination that results from AI, is already addressed by our anti-discrimination laws which are rights based and not technology specific. We are concerned that the definition of algorithmic discrimination in **AB 2930** is not sufficiently tied to those laws, such as the Fair Employment and Housing Act, or the Unruh Civil Rights Act. Rather than tying the definition to long standing statutes and jurisprudence, the bill introduces new concepts such as “contribut[ing] to unjustified differential treatment or impacts” will lead to confusion and interfere with our ability to ensure adherence to existing non-discrimination laws.

Next, we strongly believe that to accomplish this objective, the focus needs to be not on the technology itself, or the low-risk applications of it, but rather, on high-risk uses of ADT. Unfortunately, the definitions of “automated decision tool” and “consequential decision” are overly broad and at times vague, creating significant concerns that the bill reaches beyond high-risk uses. Under its definition of an automated decision tool, **AB 2930** regulates systems or services that use AI and that have been specifically developed and marketed to make, or specifically modified to make, or be a controlling factor in making, consequential decisions. Meaning, it captures even those systems where human judgment and supervision is involved and are ultimately the controlling factors in making the decision – not the ADT itself.

In turn, “consequential decision” means “a decision *or judgment* that has a legal, *material*, or *similarly significant effect* on an individual’s life relating to government benefits or services, assignments of penalties by government, or the *impact* of, or the cost, terms, or availability of *any*” of 11 different subjects ranging from employment, education, health care, family planning, housing, essential utilities, financial services, criminal justice system, legal services, voting, or access to benefits or services or assignment of penalties (emphasis added). We are concerned that various elements of the definition are too subjective, including its focus on *impact*, instead of access. Additionally, any one of the 11 identified sectors/use cases could easily be the subject of individual legislation, given the complexity of the issues involved that varies from industry to industry, yet none of them are adequately defined to make clear the types of decisions or judgments that are subject to these regulations. At the same time, qualifying certain use cases as consequential by default (e.g. use of AI for recruitment) can be too rigid, undermining the potential benefits of the AI application, and ignoring available mitigation measures.

Furthermore, where a consequential decision is made solely based on the output of an ADT, the deployer of that ADT must, “if technically feasible”, accommodate the person’s request to opt-out of the ADT and be subject to an “alternative selection process or accommodation,” without any clarity as to what is considered “technically feasible” and without any limitations on what might be considered a reasonable accommodation request. (See Proposed Section 22756.2.) Anything is *technically* feasible, with sufficient time and money, yet having standalone processes simply is not practically feasible in many, if not most, contexts as discussed further below.

These issues all become of even greater concern because the bill applies to businesses of *all* sizes that use ADTs. **AB 2930's** qualified exemption for businesses with fewer than 25 employees does not sufficiently address concerns that the bill applies to small businesses. For example, consider if a business uses a technology that can evaluate whether the business is ready to expand or needs additional employees to be hired. That presumably would be a use of ADT to make a consequential decision. Or consider if a business relies on their calendars or appointment calendaring technology to determine if they

can provide an appointment, service, or interview to an individual needing that service or job by a date certain. Have they then used ADT to make a consequential decision? If so, they would now be subject to **AB 2930** and subject to potentially severe fines and litigation costs in actions brought by the AG or other public prosecutors – fines and costs that can be crushing to all businesses, but particularly devastating to small businesses.

AB 2930's impact assessments and notice requirements raise significant issues, chief among them being that an opt-out right is wholly unrelated and unnecessary to ensuring the responsible development and deployment of ADTs.

At or before the time an ADT is used to make a consequential decision, **AB 2930** requires a deployer to provide notice to a person that is the subject of the consequential decision that an ADT is being used to make, or will be a controlling factor in making, the consequential decision, and provide each person an opportunity to opt-out of the ADT if “technically feasible”. It is critical to consider the practical impact of this requirement. First, as noted above, the focus should be on situations where human intervention is not involved, not just where ADT is “a” controlling factor among other controlling factors. Second, there are many circumstances in which providing notice and/or opportunity to opt-out is not practicable and could have significantly negative consequences on the impacted party.

By way of just one example, consider if a hospital emergency room had to stop vital treatment and notify individuals whose lives are at stake that they are deploying a triage tool to help identify the order of treatment given limited resources, and an opportunity to opt-out of that system. Simply put, this could result in dire outcomes that we do not believe this bill intends to have. Ultimately, we are concerned that the potential for steep fines will have a significant chilling effect on innovation in California and on access to important technology – including technologies that could reduce the instances of human bias and discrimination.

There are a host of other concerns identified with **AB 2930**, including in relation to its impact assessments and notice requirements. Consider the following:

- **AB 2930** requires both deployers and developers to conduct assessments, separate and apart from one another. In each of these assessments, for example, they must provide “[a] description of the automated decision tool’s outputs and how they are used to make, or be a controlling factor in making, a consequential decision.” However, the former may not have the knowledge and expertise of the underlying technology created by the latter to understand all the varying uses aside from their own intended use; and the latter may not have knowledge or control over the modification and potential uses of that modified technology by the former. Yet both are required to conduct a full assessment, including an analysis of the potential adverse impacts that could vary depending on those exact factors.
- **AB 2930** not only requires impact assessments of the same tool to be conducted annually regardless of material changes, but it also requires that they be made “as soon as feasible” with respect to any “significant update”. The former is undefined and the latter is ambiguous at best as it is defined to mean not only a new version or new release, but also any “other update” to an ADT that changes its use case, key functionality, or expected outcomes. Additionally, the annual nature of these assessments will be burdensome, particularly for smaller businesses.

AB 2930 requires additional confidentiality protections

AB 2930 rightfully includes language protecting against the disclosure of trade secrets. We appreciate the additional confidentiality protections in the bill that have been added since AB 331 but are disappointed to find that the protections remain far too narrow. Under the bill, a trade secret contained in an impact assessment disclosed to the CRD is exempt from the California Public Records Act (CPRA). It also now provides that in complying with a CPRA request, the CRD, or an entity with which an impact assessment was shared under the bill, must redact any trade secret from the impact assessment. However, there needs to be greater confidentiality protections in the bill, including not only expressly exempting these assessments from the California Public Records Act, in their entirety, but also limiting sharing among state agencies without any trigger. Insufficient confidentiality protections undermine candor and thoughtful assessments and creates concern over those assessments becoming fodder for litigation.

AB 2930 raises significant concerns around allowing enforcement by multiple entities and imposing punitive fines, with little to no opportunity to cure

We greatly appreciate the deletion of the private right of action that was previously in AB 331. Compliance is not easy to achieve in areas where massive changes in public policy are sought and where the state of law and technology are not only complicated but constantly evolving, a private right of action would have been highly problematic and chilling of innovation. That being said, businesses are still subject to civil enforcement by not only the Attorney General (AG) but by all other public attorneys (city prosecutors, district attorneys, city attorneys, and county counsel), in addition to the administrative enforcement of the Civil Rights Department (CRD). We feel it critical that the law be subject to a single enforcer. At the very least, civil enforcement should be limited to the Attorney General, particularly given the subjectivity and vagueness involved in many aspects of this bill. Otherwise, businesses can fall subject to different interpretations of the law and be found noncompliant for the same actions in one jurisdiction that would be compliant in another jurisdiction.

We also have concerns about the punitive nature of the fines imposed by **AB 2930** for each violation, and the lack of clarity around what constitutes a single “violation” for these purposes. In the context of an administrative action, a \$10,000 fine applies to each violation (compare this, for example, to California’s landmark data privacy law, which includes fines of not more than \$2,500 or, in the case of an intentional violation, not more than \$7,500). In the context of a civil action by a public prosecutor, a \$25,000 fine is applied, per violation, for algorithmic discrimination. Such fines are available even without evidence of any actual harm to the consumer. Furthermore, these fines would become even more punitive, depending on what is considered a single violation, as opposed to multiple violations. As currently drafted, an argument could be made that a violation constitutes not only a failure to complete an impact assessment altogether but also any single deficiency within that impact assessment. Even then, it is unknown whether the number of violations is based on that single error, or by any single error multiplied by the number of individuals who were *potentially* impacted by the use of a particular type of ADT for which an assessment was required or who received an inadequate notice.

And while we support the need for an opportunity to cure, conceptually, as drafted **AB 2930’s** right to cure is inadequate and illusory. Whereas the bill authorizes a court to award to any prevailing plaintiff injunctive relief, declaratory relief, reasonable attorney’s fees and litigation costs, and a civil penalty of \$25,000 per violation in an action for a violation involving algorithmic jurisdiction, the right to cure only applies to injunctive relief. Even still, no company can realistically avail themselves of this right under **AB 2930**, as it requires them to not only cure the noticed violation, but also provide the person giving the notice an express written statement that the violation has been cured *and that no further violations shall occur*. Due to the complexity and evolving nature of this technology, as well as the bill’s breadth and vagueness issues, it is unrealistic for company to be asked to sign such a statement.

AB 2930 requires preemption language to avoid significant implementation problems and disparate protections depending on the city or county a person lives in.

Unlike earlier iterations of AB 331, **AB 2930** does not include any language ensuring that no city or county can adopt, maintain, enforce, or continue in effect any law, regulation, rule, requirement, or standard related to the performance of an impact assessment or governance program, or the equivalent thereof of an impact assessment. California has 482 municipalities spread across 58 different counties. When a law is a matter of statewide concern it is critical to include preemption. (See e.g. Civ. Code Section 1798.180). AI, including issues around bias and automated decision tools, are decisions that should be made if not at the federal level, at a state one. Otherwise, imagine if even a quarter of these localities have disparate rules for impact assessments. Imagine that the conflicts are so significant that businesses must conduct multiple assessments, in addition to the ones they must conduct pursuant to state laws and even international ones. Worse yet, what if those impact assessments preclude the development or deployment of certain ADT in one city, but not the next? And what if that impacts the quality of, or even access to, services in one city to the next?

Other active legislative and regulatory efforts on related topics similarly open the door for vast confusion and increase the likelihood of conflicting rules and regulations. **AB 2930**, while specifically focused on automated decisionmaking tools, is not the only measure related to AI and issues around bias, discrimination, or other risks to consumers. Last year, in addition to AB 331, there were at least five AI bills

and two regulatory bodies grappling with the issue in one form or another. This year, there are roughly 10 times as many AI bills, which already creates enormous potential for conflict. Even more problematic: numerous departments and agencies are tasked with examining issues around AI or GenAI in addition to the two regulatory bodies working on regulations on ADTs or algorithmic discrimination.

We are especially concerned about agencies and departments getting ahead of the Legislature and Governor on matters of such statewide importance. The clearest example of this lies with the California Privacy Protection Agency. Proposition 24 of 2020 required the CPPA to issue regulations “governing access and opt-out rights related to businesses’ use of automated decision-making technology, including profiling and requiring businesses’ response to access requirements to include meaningful information about the logic involved in these decisionmaking processes, as well as a description of the likely outcome of the process with respect to the consumer.” (Civ. Code Sec. 1798.185.) Last month, the divided Agency Board voted 3-2, to advance their draft regulations, including regulations on automated decision-making tools, to formal rulemaking. As we have testified numerous times at their Board meetings, including the meeting in March, the draft regulations go beyond the authority granted by the voters and in fact veer into broader AI regulation in general, which should be addressed by this Legislature and Governor – not that Agency.

In addition to all of this, the Civil Rights Council within the Civil Rights Department (formerly, the Department of Fair Employment and Housing) has also been working on regulating the use of AI and machine learning in connection with employment decision-making. Last year, the Council published draft modifications to their proposed employment regulations regarding “automated decision systems” in their effort to incorporate such technology into existing rules regulating California employment and hiring practices, yet it is unclear when (or if) these regulations will advance into formal rulemaking.

With all these moving parts, it is difficult to foresee how such laws and regulations will layer on top of one another and whether there will be conflicting public policy around the use of such tools and technologies. While we are confident that the Legislature will review legislation for such conflicts, that does nothing to address the situation of agencies and departments doing their own rules—even ones that are arguably not supported in statutory authority. Understandably, our members are alarmed by the likelihood of conflict and confusion at the conclusion of these efforts that are being run in parallel to each other, without sufficient coordination or consideration of the other efforts underway. If **AB 2930** is passed, we hope it will include some protection against agencies and departments stepping on the Legislature and Governor’s authority.

Ultimately, overregulation in this space can easily undermine many beneficial uses of ADT—including the ability develop and deploy these tools in a manner that can in fact reduce the instances and effects of human bias. To that end, we look forward to negotiating amendments that provide greater clarity and precision to avoid unintended consequences but, for all the aforementioned reasons, must **OPPOSE AB 2930 (Bauer-Kahan)** in its current form.

Sincerely,



Ronak Daylami
Policy Advocate
on behalf of

American Council of Life Insurers
American Property Casualty Insurance Association
American Staffing Association
Association of California Life and Health Insurance Companies
Association of National Advertisers
California Association of Realtors
California Chamber of Commerce
California Hotel & Lodging Association
California Manufacturers & Technology Association
California Retailers Association
California Staffing and Recruiting Association
Chamber of Progress

Civil Justice Association of California
Computer and Communications Industry Association (CCIA)
CTIA
Insights Association
Los Angeles Area Chamber of Commerce
National Association of Mutual Insurance Companies
Personal Insurance Federation of California
State Privacy & Security Coalition
TechNet

cc: Legislative Affairs, Office of the Governor
Estefani Avila, Office of Assemblymember Bauer-Kahan
Consultant, Assembly Privacy and Consumer Protection Committee
Liz Enea, Consultant, Assembly Republican Caucus

RD:ldl